

Building An Application Using MapReduce and Workflows

Exercise Objectives

In this exercise, you will:

- Add a MapReduce job and Workflow to the example project
 - Run the MapReduce job
 - Run the Workflow
 - View operations and results in Reactor Dashboard
-

Exercise Steps

- Add imports
 - Modify `pom.xml`
 - Modify `ApplicationSpecification`
 - Modify the Procedure
 - Define the MapReduce job
 - Define `beforeSubmit`
 - Define `Mapper`
 - Define `Reducer`
 - Define the Workflow
 - Build, deploy, run and test
-

Add Imports

Add these imports:

```
import com.continuity.api.data.dataset.table.Put;
import com.continuity.api.mapreduce.AbstractMapReduce;
import com.continuity.api.mapreduce.MapReduceContext;
import com.continuity.api.mapreduce.MapReduceSpecification;
import com.continuity.api.schedule.Schedule;
import com.continuity.api.workflow.Workflow;
import com.continuity.api.workflow.WorkflowSpecification;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.MD5Hash;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
```

Modify pom.xml

```
<dependency>
  <groupId>org.apache.hadoop</groupId>
  <artifactId>hadoop-common</artifactId>
  <version>${hadoop.version}</version>
  <scope>provided</scope>
  <exclusions>
    <exclusion>
      <groupId>io.netty</groupId>
      <artifactId>netty</artifactId>
    </exclusion>
    <exclusion>
      <groupId>com.sun.jersey</groupId>
      <artifactId>jersey-server</artifactId>
    </exclusion>
    <exclusion>
      <groupId>org.jboss.netty</groupId>
      <artifactId>netty</artifactId>
    </exclusion>
  </exclusions>
</dependency>
```

Modify `ApplicationSpecification`

Replace `.noMapReduce()` with:

```
.withMapReduce()  
  .add(new SentimentAnalysisMapReduce())
```

Replace `.noWorkflow()` with:

```
.withWorkflows()  
  .add(new SentimentAnalysisWorkflow())
```

Modify `SentimentAnalysisProcedure`

Add a new handler:

```
@Handle("reductions")
public void sentimentReductions(ProcedureRequest request, ProcedureResponder response)
throws Exception {
    String sentiment = request.getArgument("sentiment");
    if (sentiment == null) {
        response.error(ProcedureResponse.Code.CLIENT_ERROR, "No sentiment sent.");
        return;
    }
    byte[] count = sentiments.get(Bytes.toBytes(sentiment), Bytes.toBytes(sentiment));
    Map<String, Long> resp = Maps.newHashMap();
    if (count == null) {
        resp.put(sentiment, 0L);
    } else {
        resp.put(sentiment, Bytes.toLong(count));
    }
    response.sendJson(ProcedureResponse.Code.SUCCESS, resp);
}
```

Add MapReduce Job

```
public static class SentimentAnalysisMapReduce extends AbstractMapReduce {

    // Annotation indicates the DataSets used in this MapReduce
    @UseDataSet("text-sentiments")
    private SimpleTimeseriesTable textSentiments;
    @UseDataSet("sentiments")
    private Table sentiments;

    private static String sentiment_arg;

    @Override
    public MapReduceSpecification configure() {
        return MapReduceSpecification.Builder.with()
            .setName("SentimentCountMapReduce")
            .setDescription("Sentiment count MapReduce job")
            .useInputDataSet("text-sentiments") // Specify the DataSet for Mapper to read.
            .useOutputDataSet("sentiments") // Specify the DataSet for Reducer to write.
            .setMapperMemoryMB(512)
            .setReducerMemoryMB(1024)
            .build();
    }
}
```

Define the `beforeSubmit`

```
@Override
public void beforeSubmit(MapReduceContext context) throws Exception {
    Job job = context.getHadoopJob();
    sentiment_arg = context.getRuntimeArguments().get("sentiment");
    if (sentiment_arg == null) {
        sentiment_arg = "positive";
    }
    LOG.info("Start of MapReduce job for sentiment \"" + sentiment_arg + "\"");

    long endTime = System.currentTimeMillis();
    long startTime = endTime - TimeUnit.MILLISECONDS.convert(1, TimeUnit.DAYS);
    context.setInput(textSentiments, textSentiments.getInput(1, Bytes.toBytes(sentiment_arg),
                                                                startTime,
                                                                endTime));

    job.setMapperClass(SentimentMapper.class);
    job.setMapOutputKeyClass(Text.class); // Sets output key of the Reducer class
    job.setMapOutputValueClass(IntWritable.class); // Sets output value of the Reducer class
    job.setReducerClass(SentimentReducer.class);
}
```

Define Mapper

```
/**
 * A Mapper that reads the sentiments from the text-sentiments
 * DataSet and creates key value pairs, where the key is the
 * sentiment and value is the occurrence of a sentence. The Mapper
 * receives a key value pair (<byte[], TimeseriesTable.Entry>)
 * from the input DataSet and outputs data in another key value
 * pair (<Text, IntWritable>) to the Reducer.
 */
public static class SentimentMapper extends Mapper<byte[], TimeseriesTable.Entry,
                                                    Text, IntWritable> {

    // The output value
    private static final IntWritable ONE = new IntWritable(1);

    @Override
    public void map(byte[] key, TimeseriesTable.Entry entry, Context context)
        throws IOException, InterruptedException {
        // Send the key value pair to Reducer.
        String sentiment = Bytes.toString(key);
        context.write(new Text(sentiment), ONE);
    }
}
```

Define Reducer

```
/**
 * Aggregates the number of sentences per sentiment and store the results in a Table.
 */
public static class SentimentReducer extends Reducer<Text, IntWritable, byte[], Put> {
    public void reduce(Text sentiment, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        long count = 0L;
        // Get the count of sentences
        for (IntWritable val : values) {
            count += val.get();
        }
        // Store aggregated results in output DataSet.
        // Each sentiment's aggregated result is stored using the sentiment as a key.
        context.write(Bytes.toBytes(sentiment.toString()),
            new Put(Bytes.toBytes(sentiment_arg),
                Bytes.toBytes(sentiment_arg), count ));
    }
}
```

Define `onFinish`

```
@Override
public void onFinish(boolean succeeded, MapReduceContext context) throws Exception {
    LOG.info("Action taken on MapReduce job for sentiment \"" + sentiment_arg + "\" : " +
        (succeeded ? "" : "un") + "successful completion");
}

} // Closes class SentimentAnalysisMapReduce
```

Define `SentimentAnalysisWorkflow`

```
/**
 * Implements a simple Workflow with one Workflow action to run
 * the SentimentAnalysisMapReduce MapReduce job with a schedule
 * that runs every day at 11:00 A.M.
 */
public class SentimentAnalysisWorkflow implements Workflow {

    @Override
    public WorkflowSpecification configure() {
        return WorkflowSpecification.Builder.with()
            .setName("SentimentAnalysisWorkflow")
            .setDescription("SentimentAnalysisWorkflow description")
            .onlyWith(new SentimentAnalysisMapReduce())
            .addSchedule(new Schedule("DailySchedule", "Run every day at 11:00 A.M.", "0 11 * * *",
                                     Schedule.Action.START))
            .build();
    }
}
```

Build, Deploy and Test

- Build using `mvn clean package`
 - Deploy the jar to Reactor after stopping any existing Flows and Procedures
 - Run the MapReduce job, using the Dashboard to set the sentiment it reduces as a Runtime Argument
 - Check the results using the Procedure's *reductions* handler, passing in the same sentiment you are checking as a parameter
 - Run the Workflow manually
-

Exercise Summary

You should now be able to:

- Add MapReduce jobs and Workflows to a project
 - Run MapReduce jobs
 - Run Workflows
 - View operations and results in the Reactor Dashboard
-

Exercise Completed

[Chapter Index](#)