

Flow Partitioning

Module Objectives

In this module, you will learn:

- Flowlet batch execution
 - Instances and scaling of Flowlets
 - Flow partitioning
 - Different strategies for partitioning
 - Configuring the number of Flowlet instances
-

Recap of the Flow System 1/2

Flows:

- User-implemented real-time stream processors
- Comprised of one or more **Flowlets** wired together into a directed acyclic graph or DAG

Flowlets:

- Pass data between one another
 - Able to perform custom logic
 - Execute data operations for each individual data object it processes
 - Processes the data objects from its input one by one
 - Multiple inputs are consumed (by default) in a round-robin fashion
-

Recap of the Flow System 1/2

When processing a single input object, all operations, including the removal of the object from the input, and emission of data to the outputs, are executed in a transaction

This provides:

- *Atomicity, Consistency, Isolation, and Durability* (ACID) properties
 - Helps assure a unique and core property of the Flow system: atomic and "exactly-once" processing of each input object by each Flowlet in the DAG
-

Flowlet Batch Execution

By default, a Flowlet processes a single data object at a time within a single transaction

Increase throughput by processing batches of data objects within the same transaction:

```
@Batch(100)
@ProcessInput
public void process(Iterator<String> words) {
    ...
}
```

In this example, up to 100 data objects can be read from the input and processed at one time

Flowlets, Instances and Scaling

- Can have one or more instances of any given Flowlet
 - Each consumes a disjoint partition of each input queue
 - Enables you to scale your application to meet capacity at runtime
 - Can control the number of instances, either:
 - Programmatically using the REST API, or via
 - The Continuuity Reactor Dashboard
 - In the **Local Reactor**, multiple Flowlet instances are run in threads, so in some cases actual performance may not be improved
 - In the **Hosted** and **Enterprise Reactors** each Flowlet instance runs in its own Java Virtual Machine (JVM) with independent compute resources
 - Scaling the number of Flowlets can improve performance and have a major impact depending on the implementation
-

Flow Partitioning Strategies

With have multiple instances of a Flowlet, the input queue is partitioned among the Flowlets

A Flowlet can specify one of these three partitioning strategies:

- **First-in first-out (FIFO):** Default strategy
 - **Round-robin:** The number of items is distributed evenly among the instances
 - **Hash-based:** The emitting Flowlet annotates each data object with a hash key
-

Details of Flow Partitioning Strategies 1/2

First-in first-out (FIFO):

- Every Flowlet instance receives the next available data object in the queue
- Since multiple consumers may compete for the same data object, access to the queue must be synchronized
- This may not always be the most efficient strategy

Round-robin:

- The number of items is distributed evenly among the instances
 - In general, the most efficient partitioning, more efficient than FIFO
 - Not ideal when the application needs to group objects into buckets according to business logic
-

Details of Flow Partitioning Strategies 2/2

Hash-based:

- The emitting Flowlet annotates each data object with a hash key
 - Ensures that all objects of a given key are received by the same consumer instance
 - Useful for aggregating by key and can help reduce write conflicts
-

Example Flow Partitioning Strategy: FIFO

A Flowlet that counts words and uses the default strategy of FIFO:

```
public class Counter extends AbstractFlowlet {

    @UseDataSet("wordCounts")
    private KeyValueTable wordCountsTable;

    @ProcessInput("wordOut")
    public void process(String word) {
        this.wordCountsTable
            .increment(Bytes.toBytes(word), 1L);
    }
}
```

Example Flow Partitioning Strategy: Round-robin 1/2

To increase throughput when this Flowlet has many instances, specify round-robin partitioning:

```
@RoundRobin
@ProcessInput("wordOut")
public void process(String word) {
    this.wordCountsTable.increment(Bytes.toBytes(word), 1L);
}
```

Example Flow Partitioning Strategy: Round-robin 2/2

If we have three instances of this Flowlet, every instance will receive every third word

For the sequence of words in the sentence

I scream, you scream, we all scream for ice cream:

- The first instance receives: *I scream scream cream*
- The second instance receives: *scream we for*
- The third instance receives: *you all ice*

Potential problem: the first two instances might both attempt to increment the counter for the word *scream* at the same time, leading to a write conflict

Example Flow Partitioning Strategy: Hash-based 1/3

To avoid conflicts, use hash-based partitioning:

```
@HashPartition("wordHash")
@ProcessInput("wordOut")
public void process(String word) {
    this.wordCountsTable.increment(Bytes.toBytes(word), 1L);
}
```

Only one of the Flowlet instances will receive the word *scream*, and there are no more write conflicts

Example Flow Partitioning Strategy: Hash-based 2/3

To use, the emitting Flowlet must annotate each data object with the partitioning key:

```
@Output("wordOut")
private OutputEmitter<String> wordOutput;
...
public void process(StreamEvent event) {
    ...
    // emit the word with the partitioning
    // key name "wordHash"
    wordOutput.emit(word, "wordHash", word.hashCode());
}
```

Example Flow Partitioning Strategy: Hash-based 3/3

- The emitter must use the same name ("wordHash") for the key that the consuming Flowlet specifies as the partitioning key
 - If the output is connected to more than one Flowlet, you can also annotate a data object with multiple hash keys
 - Each consuming Flowlet can then use different partitioning
 - Useful if you want to aggregate by multiple keys, such as counting purchases by product ID as well as by customer ID
-

Combining Flow Partitioning and Batch Execution

Partitioning can be combined with batch execution:

```
@Batch(100)
@HashPartition("wordHash")
@ProcessInput("wordOut")
public void process(Iterator<String> words) {
    ...
}
```

Number of Flowlet Instances: REST API 1/3

You can query and set the number of instances executing a given Flowlet by using the `instances` parameter with HTTP GET and PUT methods:

```
GET <base-url>/apps/<app-id>/flows/<flow-id>/flowlets/  
    <flowlet-id>/instances  
PUT <base-url>/apps/<app-id>/flows/<flow-id>/flowlets/  
    <flowlet-id>/instances
```

with the arguments as a JSON string in the body:

```
{ "instances" : <quantity> }
```

Parameter : *Description*

<app-id> : Name of the Application being called <flow-id> : Name of the Flow <flowlet-id> : Name of the Flowlet <quantity> : Number of instances to be used

Number of Flowlet Instances: REST API 2/3

Example

Find out the number of instances of the Flowlet *saver* in the Flow *WhoFlow* of the Application *HelloWorld*

```
GET <base-url>/apps/HelloWorld/flows/WhoFlow/flowlets/  
saver/instances
```

Number of Flowlet Instances: REST API 3/3

Example

Set the number of instances of the Flowlet *saver* in the Flow *WhoFlow* of the Application *HelloWorld*

```
PUT <base-url>/apps/HelloWorld/flows/WhoFlow/flowlets/  
saver/instances
```

with the arguments as a JSON string in the body:

```
{ "instances" : 2 }
```

Setting Number of Flowlet Instances: Dashboard 1/2

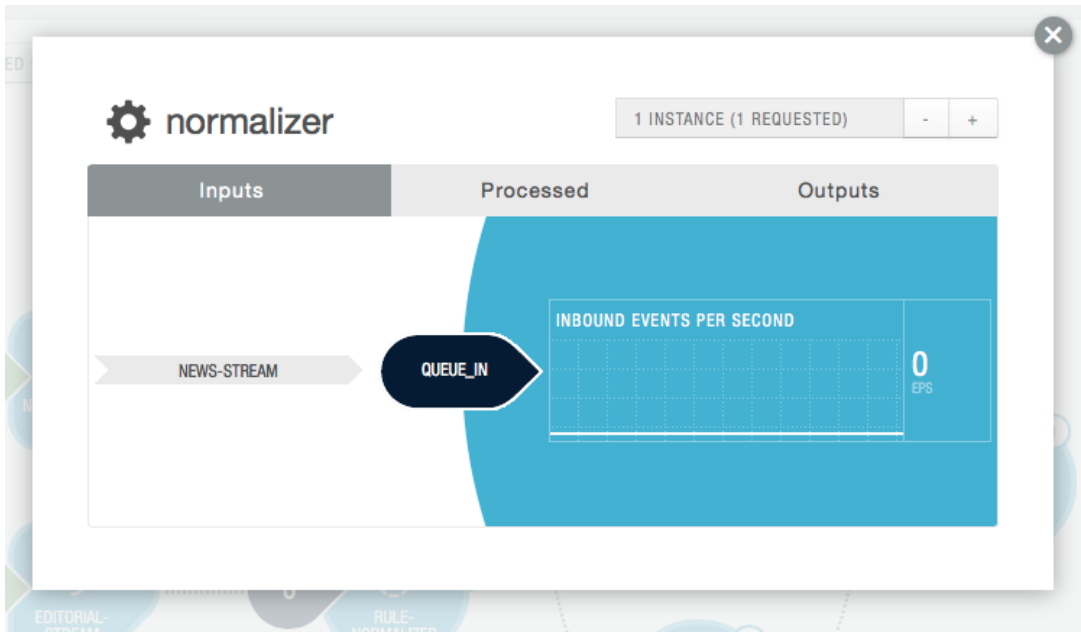
A Flowlet icon in the DAG shows:

- The name of the Flowlet
- The number of events processed in the current sampling period
- The number of instances of that Flowlet in the small circle in the upper right of the icon:



Setting Number of Flowlet Instances: Dashboard 2/2

Clicking on a Flowlet’s icon in the DAG brings up the configuration dialog for setting the number of instance of the Flowlet:



Module Summary

You should now be able to:

- Specify Flowlet batch execution
 - Describe the different Flow partitioning strategies
 - Select and specify an appropriate Flow partitioning strategy
 - Configure the number of Flowlet instances
-

Module Completed